

TODOs

■ write c1Introduction	1
■ add illustration for introduction	1
■ add citation	1
■ write c2deterministic AUC scores, compared to Mirowski AUC scores. And number of patients with IoC.	2
■ add citation	3
■ write comparison between c2deterministic and Mirowski et al. work	12
■ write about probabilistic forecast skill evaluation methods	18
■ insert illustration of data collection	20
■ write about train-test split (offline-online modes)	20
■ fix class labels in legend, remove redundant axis label	21
■ remake seperability test figures	22
■ reconsider if novelty score is necessary	22
■ write caption	23
■ replace with vectorized image	24
■ add original sample	24
■ remake example likelihood estimation figures	25
■ remake circadian profile diagrams panel horizontal	25
■ format as definition	25
■ write list of ways this work differs from prior works	27
■ write about Cox process (a.k.a. Inhomogeneous Poisson processes)	27
■ make figure of inferring latent intensity	27
■ add citation	28
■ check which threshold made the roc curve fig and write it	32
■ rewrite Linear probing section	34
■ add citation	36

■ add citation	36
■ report Brier Score	37
■ report Snyder 2008	37
■ write about model validation	38
■ rewrite discussion. Points to include: (1) future work on hierarchical patient modeling	38

Exploring the Science of Seizure Timing: from Deterministic to Probabilistic Forecasts

*Thesis submitted in partial fulfillment of the requirements for
the degree of “MASTER OF SCIENCE”*

By
NOAM SIEGEL



Submitted to the Department of Computer Sciences at
Ben-Gurion University of the Negev

JULY 2022

BEER SHEVA

Exploring the Science of Seizure Timing: from Deterministic to Probabilistic Forecasts

*Thesis submitted in partial fulfillment of the requirements for
the degree of "MASTER OF SCIENCE"*

By
NOAM SIEGEL

This work was carried under the supervision of

DR. OREN SHRIKI
AND
DR. DAVID TOLPIN

The Department of Computer Sciences

Approved by the advisors: Oren Shriki

Approved by TODO:

JULY 2022

BEER SHEVA

Acknowledgements

First and foremost, I would like to thank...

To

"The theory of probabilities is at bottom nothing but common sense reduced to calculus"

Essai Philosophique sur les Probabilités
Pierre-Simon Laplace

Abstract

My abstract

Contents

Abstract	ii
1. Introduction	1
2. Deterministic Seizure Prediction	2
2.1. Chapter Overview	2
2.1.1. Evaluation Metrics	2
2.1.2. Epilepsy	3
2.1.3. EEGs: electric potential recordings	3
2.1.4. Data	3
2.2. Methods	4
2.2.1. Preprocessing	4
2.2.2. Binary Labelling for Seizure Prediction	5
2.2.3. Bivariate Feature Extraction Methods	5
2.2.4. Classifiers	8
2.3. Results	8
2.3.1. Classifier and Feature Comparison	9
2.3.2. Training Efficiency	9
2.3.3. Results for Best Classifiers	11
2.4. Discussion	11
2.5. Conclusion	11
3. Bayesian Seizure Likelihood Estimation	17
3.1. Chapter Overview	17
3.1.1. General approach	17
3.1.2. Problem setup	17
3.1.3. Methods for evaluating forecast skill	18
3.2. Unsupervised Seizure Likelihood Estimation	18
3.2.1. Gaussian processes	18
3.2.2. Data	19
3.2.3. Embedding of EEG	20
3.2.4. Density estimation	20
3.2.5. Anomaly detection and ictality	20
3.3. Weakly Supervised Bayesian Seizure Likelihood Estimation	24
3.3.1. Subject-dependent circadian-profile prior	25
3.3.2. The Cox process prior	27
3.3.3. Probabilistic model	28

3.3.4. Empirical Evaluation	34
4. Conclusion	36
4.1. Research questions	36
4.2. Evaluation plan	36
5. Discussion	38
5.1. Model validation	38
Bibliography	40
Appendix A. GP embedding: training details	44
Hebrew Abstract	46

1. Introduction

NS: write c1Introduction

NS: add illustration for introduction

People who suffer from epilepsy undergo recurring *seizures* - episodes that include loss of consciousness, loss of motor control, and unusual paroxysmal neuronal firing. These episodes are also termed *ictal* episodes. The uncertainty associated with seizure occurrences is deemed to be the leading cause of fear, stress and other comorbidities in patients with epilepsy. Therefore, a reliable method for estimating the likelihood of a near-term seizure is in the interest of epilepsy patients and their caregivers.

We tackle the problem of inferring the probability of an ictal episode at a given time, *conditioned* on observing a window of real-time EEG data. To the best of our knowledge, most machine-learning solutions take a fully supervised approach, therefore depending on labeled datasets for learning the discrimination function. This increases the costly dependence on expert-level domain knowledge. Even clinical-grade annotations, considered to be the "gold standard" in seizure documentation, suffer from observation bias.

To overcome this problem, we propose a weakly-supervised Bayesian framework for likelihood estimation of epileptic seizures. We introduce a multilevel probabilistic model of seizure occurrences and clinical annotations (dataset labels). The model takes into account the possibility of missed seizures unrecorded by the annotator. We implemented an algorithm to compute the inference, namely Bayesian Seizure Likelihood Estimator (BSLE). Lastly, we validate the model and examine it's applicability to a practical seizure warning system.

NS

add citation

2. Deterministic Seizure Prediction

It is accepted by the epilepsy community that predicting the oncoming of a seizure is a worthy goal [Kelley, Jacobs, and Lowenstein, 2009, Dumanis, French, Bernard, Worrell, and Fureman, 2017]. The outcome of this chapter is the implementation and evaluation of an end-to-end pattern recognition system for EEG data from patients with epilepsy.

2.1. Chapter Overview

The paper which inspired this chapter is [Mirowski, Madhavan, LeCun, and Kuzniecky, 2009]. They show one of the first works to use convolutional neural networks, the same components that boosted the success of deep-learning in image processing. In this case, the classifiers were trained on patterns created by hand-engineered features from segments of 5 minute-long EEG recordings (for example, see figure 2.5), to distinguish between preictal and interictal patterns.

While Mirowski et al. [2009] trained several classifiers (logistic regression, SVM and LeNet5-CNN), on a set of synchronicity-based features, we trained a wider variety of machine learning models, using reusable, modular, scalable python code. We performed 5-fold cross validation of 9 types of classifiers on each of 5 nonlinear and bivariate feature datasets, for a total of 675 model fittings over 3 patients.

NS: write c2deterministic AUC scores, compared to Mirowski AUC scores. And number of patients with IoC.

2.1.1. Evaluation Metrics

In this project, we report the following metrics for all features-classifiers pairs:

1. ROC AUC
2. Precision
3. Recall
4. fit time
5. score time

And for some select feature-classifier pairs we report the ROC curve as well.

2.1.2. Epilepsy

Epilepsy is characterized by pathological electric activity occurring in the brain. In some individuals, this pathology can vary with time; epilepsy can form in a nonepileptic brain, exhibit changing etiologies, and for some patients be resolved, allowing them seizure-freedom [Kandel, Schwartz, Jessell, Siegelbaum, Hudspeth, and Mack, 2000].

The classification of epilepsy into types is an ongoing debate. The epilepsies are grouped hierarchically according to seizure types, epilepsy types and epilepsy syndromes, and longitudinally according to etiologies (see Figure I. in [Scheffer, Berkovic, Capovilla, Connolly, French, Guilhoto, Hirsch, Jain, Mathern, Moshé, et al., 2017]).

2.1.3. EEGs: electric potential recordings

Electroencephalography (EEG) is a form of neuroimaging which is capable of sensing local-field-potentials near cortical neural populations, electrical activity in the brain. The EEG is an effective tool in epilepsy seizure monitoring and understanding because of its sensitivity to macroscopic neural network dynamics, such as the balance of excitatory and inhibitory processes, or measures of synchronicity among different sensor locations .

The various terms local field potentials (LFPs), electrocorticography¹ (ECoG) and electroencephalography (EEG) all refer to measurements of electric potential: in nerve tissue, on the directly exposed surface of the brain, or on the surface of the scalp, respectively. These types of recording apparatuses measure electric potential which originates from summed electric activity of populations of individual cells (e.g., neurons).

Traditionally, video-electroencephalograms (vEEG) are the central diagnostic tool used by professionals to assess epilepsy in patients. The vEEG records simultaneously both optical video information and EEG data, which allows clinicians to observe simultaneously the subject and EEG dynamics. Due to the independence between the modalities, EEG is often prevalent without video surveillance. In these cases, diagnostics rely more heavily on EEG analysis.

2.1.4. Data

EEGs may be intracranial or placed noninvasively on the scalp's surface. In any case, the resulting data is of the matrical form

$$X \in \mathbb{R}^{c \times T \cdot \text{freq}} \quad (2.1)$$

where c is the number of channels, freq is the sampling frequency measured in Hz and T is the time length of the recording measured in s (see figure ??). The entries of X are typically measured in $\mu V s$. iEEG has a higher signal-to-noise ratio, but sEEG has the advantage of noninvasiveness.

¹also termed intracranial EEG.

NS

add citation

Data collection

Typical publicly available epilepsy-seizure-prediction datasets consists of long-term raw EEG recordings, supplemented with annotations provided by approved experts [Handa, Mathur, and Goel, 2021]. In these datasets, the occurrence of seizures is reported as a list of ictal (seizure) intervals

$$I := \{I_{ictal_i}\} = \{\langle t_{start}^i | S_i | t_{end}^i \rangle\} \quad (2.2)$$

where t_{start} is the seizure onset, t_{end} the seizure offset, and S_i , when available, are the additional details reported such as seizure etiology. In this project, we used data from the Epilepsiae Dataset [Ihle, Feldwisch-Drentrup, Teixeira, Witon, Schelter, Timmer, and Schulze-Bonhage, 2012]. Table 2.1 lists raw recording length, sample frequency, number of surface electrodes and size-on-disk per patient.

Table 2.1. Data used in this chapter.

Patient Name	Length	Frequency	Electrodes	SOD
pat_3500	92.8 h	1024 Hz	32	21 GB
pat_3700	79.8 h	512 Hz	32	8.8 GB
pat_7200	94.6 h	1024 Hz	29	21 GB

2.2. Methods

In the paper which inspired this chapter, Mirowski et al. [2009] trained three variants of a neural network, namely support vector machines, logistic regression and convolutional neural network. They evaluated the classifiers on 6 types of bivariate features from the fields of correlation statistics, nonlinear signal analysis, wavelet spectral analysis, and chaos theory. Of these, we selected 5 features (cf. section 2.2.3) which cover all of the feature-category types, and used the implementations in the *mne-features* module given by [Schiratti, Le Douget, Le van Quyen, Essid, and Gramfort, 2018].

2.2.1. Preprocessing

Although it is common to apply spectral filters to EEG data to reduce line noise etc., we chose to focus on the effect of classifier and features choice on predictive performance, and thus kept preprocessing to a minimum. First, EEG channel selection was performed to select a spatially far-reaching distributed subset of 19 channels, common to all patients:

['Fp1', 'Fp2', 'F3', 'F4', 'C3', 'C4', 'P3', 'P4', 'O1', 'O2', 'F7', 'F8', 'Fz', 'Cz', 'Pz', 'T7', 'T8', 'P7', 'P8'] .

This was done to reduce data dimension as well as standardization amongst patients. The raw data was resampled to 256 Hz to reduce disk space and processing time. It was then shifted and scaled to zero mean and unit standard deviation, per patient.

2.2.2. Binary Labelling for Seizure Prediction

A practical way to characterize time-varying brain dynamics is through a sliding window approach (see Fig. 1. in [Lehnertz, Geier, Rings, and Stahn, 2017]). An EEG time series matrix X (see eq. 2.1) is partitioned temporally into non-overlapping windows x_1, \dots, x_N . Each window x_i is given a label y_i , to form a dataset $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$. Optionally, each window is transformed via a feature extraction function $f(x)$, which yields a feature-formed dataset $D_f = \{(f(x_1), y_1), \dots, (f(x_N), y_N)\}$.

In binary classification for seizure prediction, each window is matched against the seizure intervals (see eq. 2.2) to create the appropriate label:

$$y_i := \begin{cases} 0 & \text{if } preictal(x_i) \\ 1 & \text{if } interictal(x_i) \end{cases} \quad (2.3)$$

Here, $ictal(x_i)$ means that x_i temporally overlaps with some ictal interval I_j . Similarly, $preictal(x_i)$ means that x_i overlaps with a specific τ_p time interval preceding an ictal interval I_j ². Finally, $interictal(x_i)$ means that x_i does not closely precede or proceed any known seizure interval, with thresholds τ_a before and τ_b after, respectively³.

2.2.3. Bivariate Feature Extraction Methods

The deterministic classification method relies on hand-crafted, manually engineered, features. Specifically, following [Mirowski et al., 2009], we focus on bivariate measures of synchronicity between pairs of EEG channels (see figure 2.1 for examples).

For each patient, the recording is segmented into 5 minute windows. Each window is segmented into 60 frames, each 5 seconds long. Each 5 second frame is reduced into a vector of length $c \cdot (c - 1)/2$ (where c is number of channels), via one of the feature extraction methods described below. Each 5 minute window is regarded a single pattern with a single label (preictal vs. interictal).

Maximal Linear Cross Correlation (figure 2.2)

In order to quantify the similarity of two signals $\{x_i\}$ and $\{y_i\}$ the maximum of a normalized cross-correlation function is taken as a measure for lag synchronization [Rosenblum, Pikovsky, and Kurths, 1997]:

$$C_{max} = \max_{\tau} \left| \frac{C_{xy}(\tau)}{\sqrt{C_{xx}(0) \cdot C_{yy}(0)}} \right| \quad (2.4)$$

²In this project, τ_p is taken to be 1 hour.

³ τ_a and τ_b are each taken to be 4 hours.

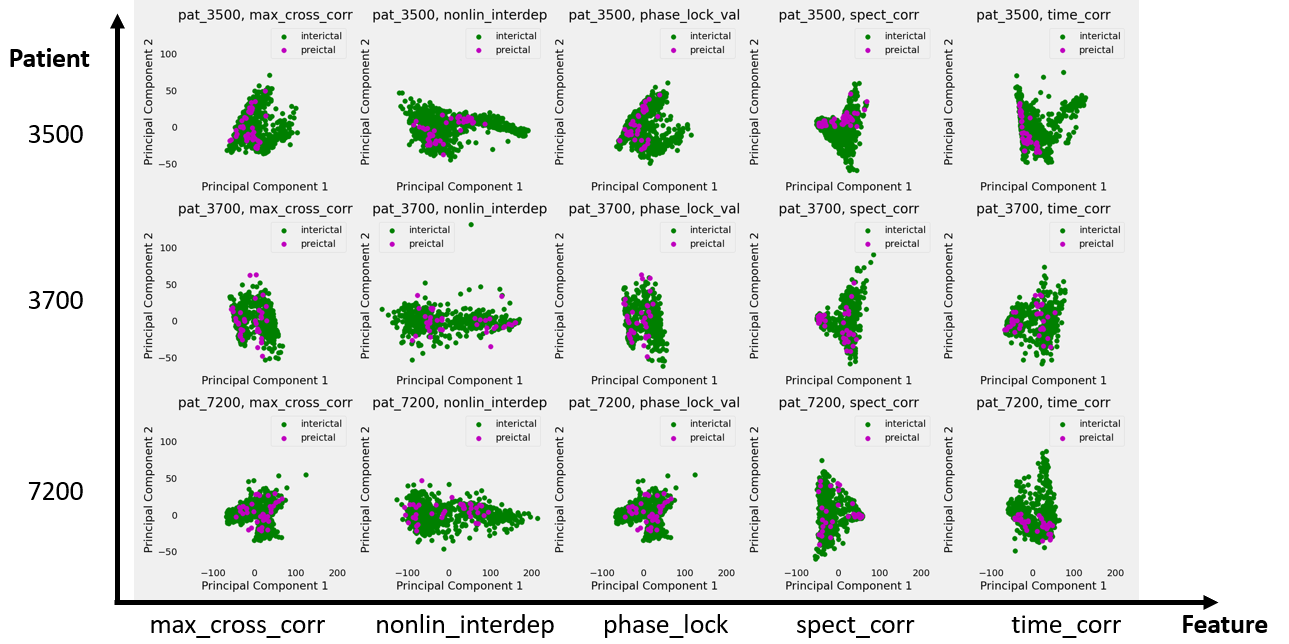


Figure 2.1. Extracted Features

Transformed datasets are visualized in the reduced 2-dimensional PCA space. Each point represents a 5-minute window (i.e., a pattern), and its color denotes class label (interictal or preictal).

Where

$$C_{xy} = \begin{cases} \frac{1}{N-\tau} \sum_{i=1}^{N-\tau} x_{i+\tau} y_i, & \text{for } \tau \geq 0 \\ C_{yx}(-\tau), & \text{for } \tau < 0 \end{cases} \quad (2.5)$$

is the linear cross-correlation function. C_{max} is confined to the interval $[0, 1]$ with high values indicating that the two signals have a similar course in time (though possibly shifted by a time lag τ) while dissimilar signals will result in values close to zero.

Phase Locking Value (figure 2.3)

Introduced in [Lachaux, Rodriguez, Martinerie, and Varela, 1999], the *Phase Locking Value* measures synchronicity between eeg channels in different locations in the brain. First, for each channel i , the instantaneous phase $\sigma_i^a(t)$ of the analytical signal $x_i^a(t)$ is extracted. Then, for each pair (i, j) of channels, we compute the modulus of the time averaged phase difference mapped onto the unit circle:

$$PLV_{ij} = \left| \frac{1}{T} \sum_t e^{i(\phi_i^a(t) - \phi_j^a(t))} \right| \quad (2.6)$$




Figure 2.2. Maximal Linear Cross-Correlation

137 **Correlation Coefficients and Eigenvalues in the Time and Frequency Domains (figure**
 138 **2.4)**

139 We compute the correlation coefficients between each pair of EEG channels, along with the eigenvalues of
 140 the correlation coefficients matrix, in both the time and frequency domains. This provides two more sets of
 141 measures of synchronicity across EEG channels.

142 **Nonlinear Interdependence (figure 2.5)**

143 The *non-linear interdependence* measure for generalized synchronization between two EEG signals $\{x_i\}$ and
 144 $\{y_i\}$ is presented in [Mormann, Andrzejak, Elger, and Lehnertz, 2007]. First, the two signals are represented
 145 as trajectories in a state space, via time-delay embedding. Then, an asymmetric statistic measuring the
 146 Euclidean distance, in reconstructed state-space, between trajectories $\{\vec{x}_i\}$ and $\{\vec{y}_i\}$ is calculated. See
 147 [Mirowski, LeCun, Madhavan, and Kuzniecky, 2008] for more details.




Figure 2.3. Phase Locking Value

2.2.4. Classifiers

We trained and tested 9 of classifiers implemented in the *scikit-learn* [Buitinck, Louppe, Blondel, Pedregosa, Mueller, Grisel, Niculae, Prettenhofer, Gramfort, Grobler, Layton, VanderPlas, Joly, Holt, and Varoquaux, 2013] ML toolkit (v1.0.1), listed in table 2.2. We chose a variety of classifiers from different families (i.e., neural networks, ensembles and decision trees).

2.3. Results

A total of 117 classifier-datasets pairs were trained and evaluated with 5-fold cross validation, yielding a total of 585 model fits.




Figure 2.4. Inter-channel correlation coefficients, with eigenvalues appended at the bottom.

2.3.1. Classifier and Feature Comparison

Figures 2.6 and 2.7 show the mean and standard-deviation for each classifier on each feature set, for each patient. It is found that the Linear SVM and Neural Net (MLPClassifier) are the top two performers consistently, sometimes followed closely by the Nearest Neighbor Classifier.

2.3.2. Training Efficiency

Figure 2.8 shows the mean time it took to fit each classifier to each feature-set, for each patient, over 5-fold CV. It is shown that the neural network, ensemble method (AdaBoost) and support-vector-machines consistently take longer than the decision tree, naive-Bayes, quadratic discriminant analysis, and random forest classifiers.




Figure 2.5. Nonlinear Interdependence

Table 2.2. Classifiers used in chapter

Classifier Name	Parameters
Nearest Neighbors	K=3
Linear SVM	kernel="linear", C=0.025
RBF SVM	gamma=2, C=1
Decision Tree	max_depth=5
Random Forest	max_depth=5, n_estimators=10
Neural Net	alpha=1, max_iter=1000
AdaBoost	default
Naive Bayes	default
QDA	default

2.3.3. Results for Best Classifiers

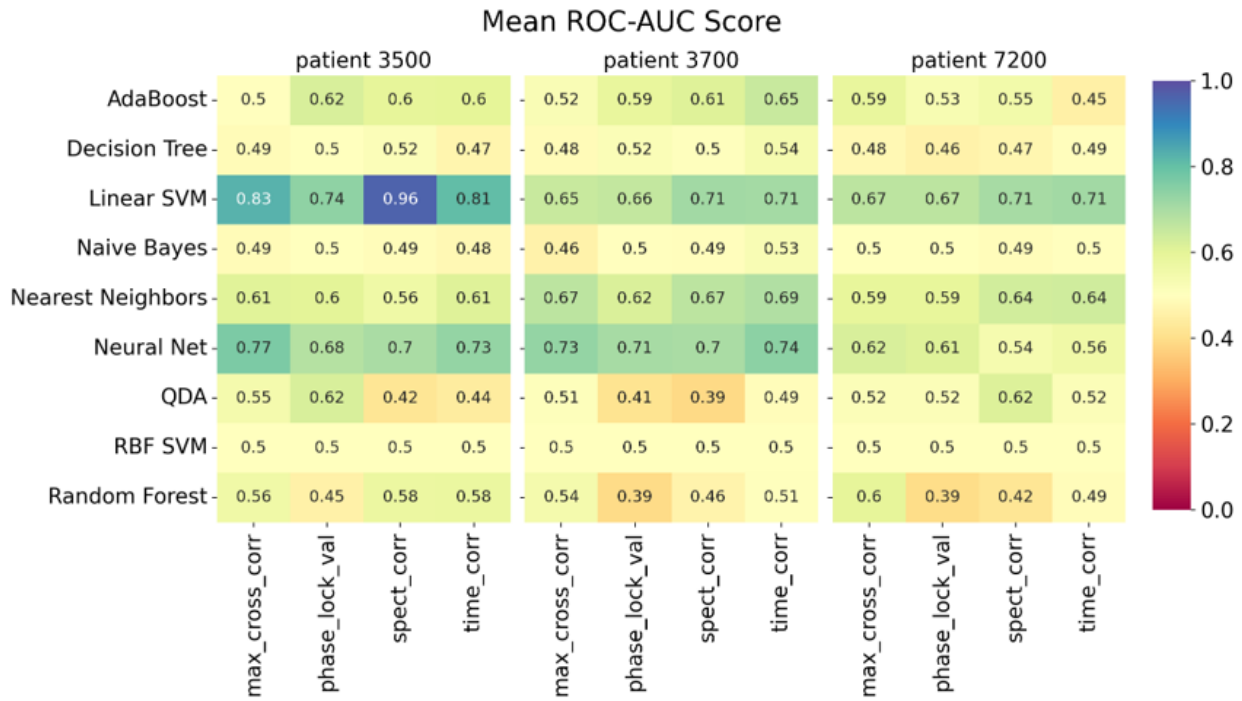

Since the Linear SVM performed the best in the previous experiments, we compared the results from the Linear SVM classifier on different feature sets. Notably, the classifier performed very similarly on all feature sets, with `spect_corr` slightly outperforming with respect to the ROC AUC score. See figures 2.9, 2.10 and 2.11.

2.4. Discussion

Mirowski et. al. evaluated their pipeline on 6-channel, intracranial data from 21 epilepsy patients. In this project, we use 19-channel data from surface EEG, from 3 patients. They reported 71% sensitivity with 0 false alarms, on 15 out of the 21 patients they assessed. We trained 117 classifiers in different combinations of classifier-dataset. For each patient, the ROC curves for the highest-roc-auc-scoring classifier are presented in figures 2.12, 2.13, and 2.14. As can be seen, only Patient 3700 achieves over 70% sensitivity at the 0 false alarms threshold. For Patients 3500 and 7200, the sensitivities at 0 false alarms are 42% and 40%, respectively. In my opinion, the leading factor in explaining this performance difference is the sensor location: intracranial in the original paper, and extracranial (surface) in our project. Another difference is the number of channels. Perhaps with the increase in channel pairs from 15 pairs to 171 pairs, the sample space grows significantly such that much more data is needed.

2.5. Conclusion

The problem of quantifying the likelihood of a seizure occurrence, namely seizure susceptibility, is an open challenge in the epilepsy research community [?]. The aim of this project is to showcase the predictive power of different classifiers and synchronicity-based features. We trained 117 classifiers in different combinations

**Figure 2.6.** Mean AUC-ROC scores for 5-fold cross validation.**Figure 2.7.** Standard deviation of AUC-ROC scores for 5-fold cross validation.

185 of classifier-datasets. Although some classifiers performed better than others, we did not find any of the
 186 feature sets to be remarkably better than the others, especially when generalizing across all 3 patients.

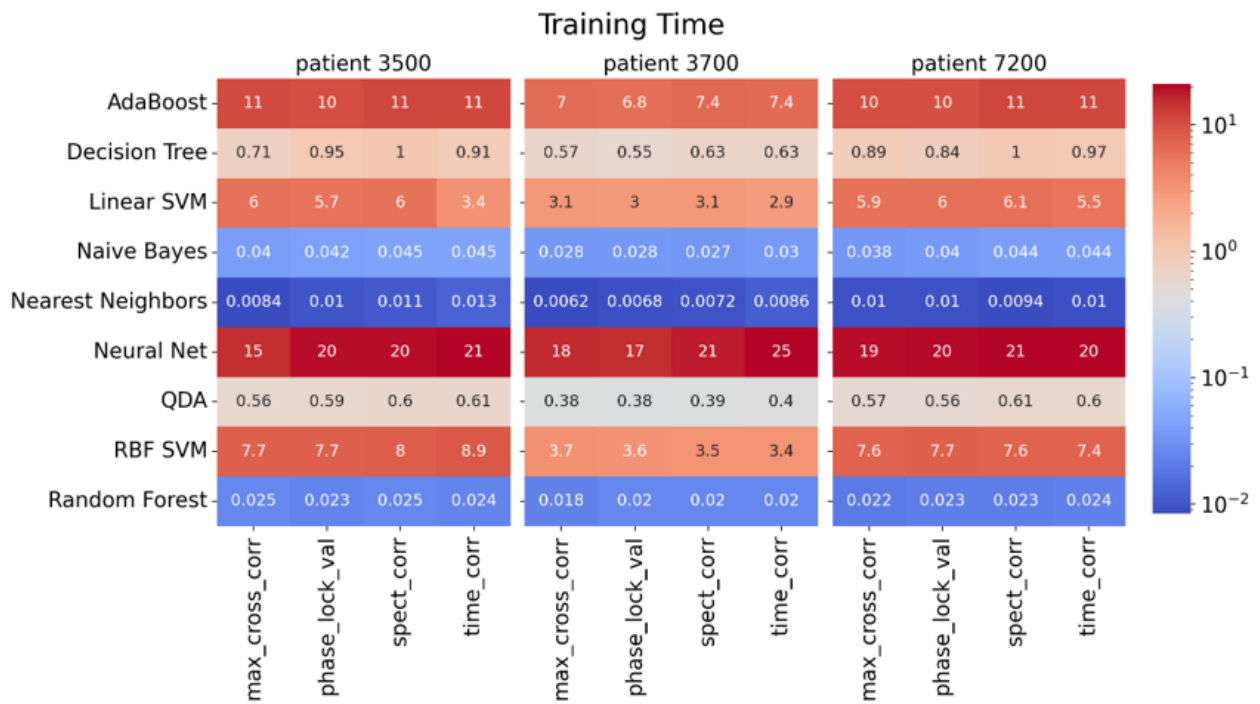


Figure 2.8. Mean training times (in seconds) for fitting the classifiers to the datasets.

NS: write comparison between c2deterministic and Mirowski et al. work




Figure 2.9. Comparison of predictive performance of Linear SVM for Patient 3500 on different feature datasets




Figure 2.10. Comparison of predictive performance of Linear SVM for Patient 3700 on different feature datasets




Figure 2.11. Comparison of predictive performance of Linear SVM for Patient 7200 on different feature datasets




Figure 2.12. Plot of Confusion Matrix and ROC Curve for the highest scoring classifier for Patient 3500




Figure 2.13. Plot of Confusion Matrix and ROC Curve for the highest scoring classifier for Patient 3700

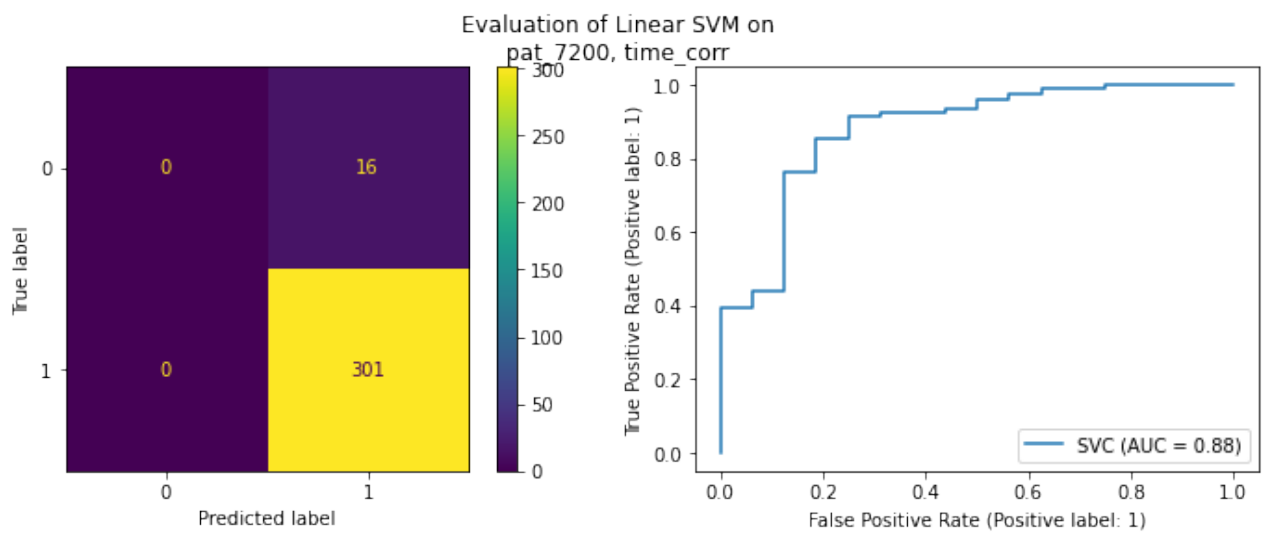


Figure 2.14. Plot of Confusion Matrix and ROC Curve for the highest scoring classifier for Patient 7200

188 3. Bayesian Seizure Likelihood 189 Estimation

190 Following a line of recent works [Karoly et al., 2017, Baud et al., 2018, Baud, Proix, Rao, and Schindler,
191 2020], we consider a probabilistic approach to seizure prediction, expressed via an algorithm we term
192 Bayesian Seizure Likelihood Estimation (BSLE)¹. This chapter focuses on the introduction of preliminary
193 concepts necessary to understand the general approach, as well as the utilities chosen in our implementation.

194 3.1. Chapter Overview

195 3.1.1. General approach

196 Our approach utilizes pure Bayesian statistics wherever possible. We develop a probabilistic model for an
197 observable EEG time series, and then incorporate weak supervision through the use of a domain-specific
198 prior. Our algorithm, Bayesian Seizure Likelihood Estimation (BSLE), performs inference computations
199 with the developed model, to assess the likelihood of an upcoming seizure. The output is tested and evaluated
200 for model fit and predictive capabilities.

201 3.1.2. Problem setup

202 Consider the problem of modeling the relationship between EEG signals (E) from the epileptic brain,
203 the occurrence of seizures (S), and an external stream of timestamps denoting seizure events, commonly
204 provided by expert annotators (A), throughout time (T).

205 Specifically, we are given a dataset $D = \{e_t, a_t\}_{t_0}^{t_f}$ where $e_t \in \mathbb{R}^{c \times \tau}$ is the observed EEG segment with
206 c -channels of duration τ recorded at time t , and $a_t \in \{0, 1\}$ is a clinically-approved annotation denoting
207 whether or not a seizure event began within the time-window $[t, t + \tau]$.

208 We wish to construct a model for $\mathbb{P}[S, E, t]$, and then apply it with Bayes' rule to infer the likelihood of a
209 seizure at time t determined by an EEG recording:

$$210 \quad \text{probability}[\text{seizure}=S \mid \text{time}=t, \text{EEG}=E] \propto \mathbb{P}[E \mid S, t] \mathbb{P}[S \mid t] \quad (3.1)$$

211 It should be evident that this procedure is general in that each component on the r.h.s. can be estimated
212 independently, and then combined via multiplication.

¹code can be found at <https://github.com/noamsgl/msc> under the MIT license.

3.1.3. Methods for evaluating forecast skill

NS: write about probabilistic forecast skill evaluation methods

3.2. Unsupervised Seizure Likelihood Estimation

3.2.1. Gaussian processes

A Gaussian process (GP) [Rasmussen, 2006] is a collection of random variables, any finite number of which have a joint Gaussian distribution.

A Gaussian process $f(x)$ is fully specified by a mean function $\mu(x)$ and a covariance function, or a kernel, $k(x, x')$, by-way-of:

$$m(x) = \mathbb{E}[f(x)] \quad (3.2)$$

$$k(x, x') = \mathbb{E}[(f(x) - m(x))(f(x') - m(x')))] \quad (3.3)$$

And it is denoted:

$$f(x) \sim \mathcal{GP}(m(x), k(x, x')) \quad (3.4)$$

In this work we will take the mean function to be zero.

parameter estimation (inference)

Gaussian processes are commonly used for time series modeling with machine learning. To see why this makes sense, imagine the input x is the time point, and the output $f(x)$ is the time series value at time x . Computationally, this is made feasible by evaluating the function's values at a finite number of points of interest. Using optimization techniques, the model's hyperparameters are inferred to match observed data by maximizing the likelihood function $p(f(x) \mid \vec{\theta})$ (termed maximum likelihood estimation, or MLE). The learned hyperparameters capture global evolutionary dynamics of the time series (see figure 3.4 in the Methods section).

The Matérn class of covariance functions

The Matérn class of covariance functions is given by:

$$k_{\text{Matern}}(x, x') = \frac{2^{(1-\nu)}}{\Gamma(\nu)} (\sqrt{2\nu}d)^\nu K_\nu(\sqrt{2\nu}d) \quad (3.5)$$

Where:

- $d = (x - x')^T \Phi^{-2} (x - x')$ is the distance between x and x' scaled by the *lengthscale* parameter Φ .
- ν is a smoothness parameter. In this work, it is taken to be $\frac{3}{2}$.
- K_ν is a modified Bessel function.

Multitask Gaussian processes

In case $f(x)$ is a vector function, multiple output functions are modeled in conjunction for the same input values, so-called multitask Gaussian process modeling. In this case, given inputs x and x' , and tasks i and j , the covariance between two datapoints and two tasks is given by:

$$k([x, i], [x', j]) = k_{\text{inputs}}(x, x') \cdot k_{\text{tasks}}(i, j) \quad (3.6)$$

Where k_{inputs} is a standard kernel (e.g., Matérn) that operates on the inputs, and k_{tasks} is a lookup table containing inter-task covariance. This is akin to capturing the inter-channel synchronicity with manually engineered features (see ??).

Gaussian mixture models [Theodoridis, 2015] are used to model the distribution of an unknown set of vectors $\{x\} \subseteq \mathbb{R}^l$ as a linear combination (i.e., a mixture) of different Gaussian distributions, that is,

$$p(x) = \sum_{k=1}^K p_k p(x | k; \zeta_k) \quad (3.7)$$

where $\{\zeta_k\}$ parametrize the individual Gaussian distributions:

$$p(x | k; \zeta_k) = p(x | k; \mu_k, \sigma_k) = \mathcal{N}(x | \mu_k, \sigma_k) \quad (3.8)$$

Fitting the model provides an approximation $\hat{p}(x | k; \zeta_k)$ of the dataset's underlying pdf, which is an estimate of the data-distribution of the GP-hyperparameters $P(E | D)$

3.2.2. Data

In this work we use the *Canine-Epilepsy Dataset* [Davis et al., 2011, UPenn and Mayo Clinic and Kaggle, 2014]. Three dogs with naturally occurring epilepsy are monitored continuously, for 330/451/475 days

for Dog 1, Dog 2, and Dog 3, respectively. The monitoring systems consist of 16-electrode EEG sensors implanted in each brain, digitally sampling at the rate of 400 Hz. In addition, a sequence of timestamps are provided for each dog, marking the observed seizure events (a.k.a. annotations).

NS: insert illustration of data collection

Train-test split

NS: write about train-test split (offline-online modes)

3.2.3. Embedding of EEG

Our unsupervised embedding method appears to show high distinguishability between interictal and ictal data. As seen in figure 3.1, in the GP parameter space the seizures are well separated from the non-seizure segments.

In order to further demonstrate the distinguishability between the interictal and ictal states, we show in figure 3.2 that a linear-kernel SVM fit to the training set achieves 0.91 AUC on the test set in the single-channel case. The double-channel case scores higher.

3.2.4. Density estimation

We use density-estimation of the interictal segments to contain the model of normal EEG. The data likelihood was maximized for a Gaussian mixture model with 2 components, fit to the interictal segments. Figure 3.8 shows the the contour plot of the log-likelihoods scored by the GMM.

3.2.5. Anomaly detection and ictality

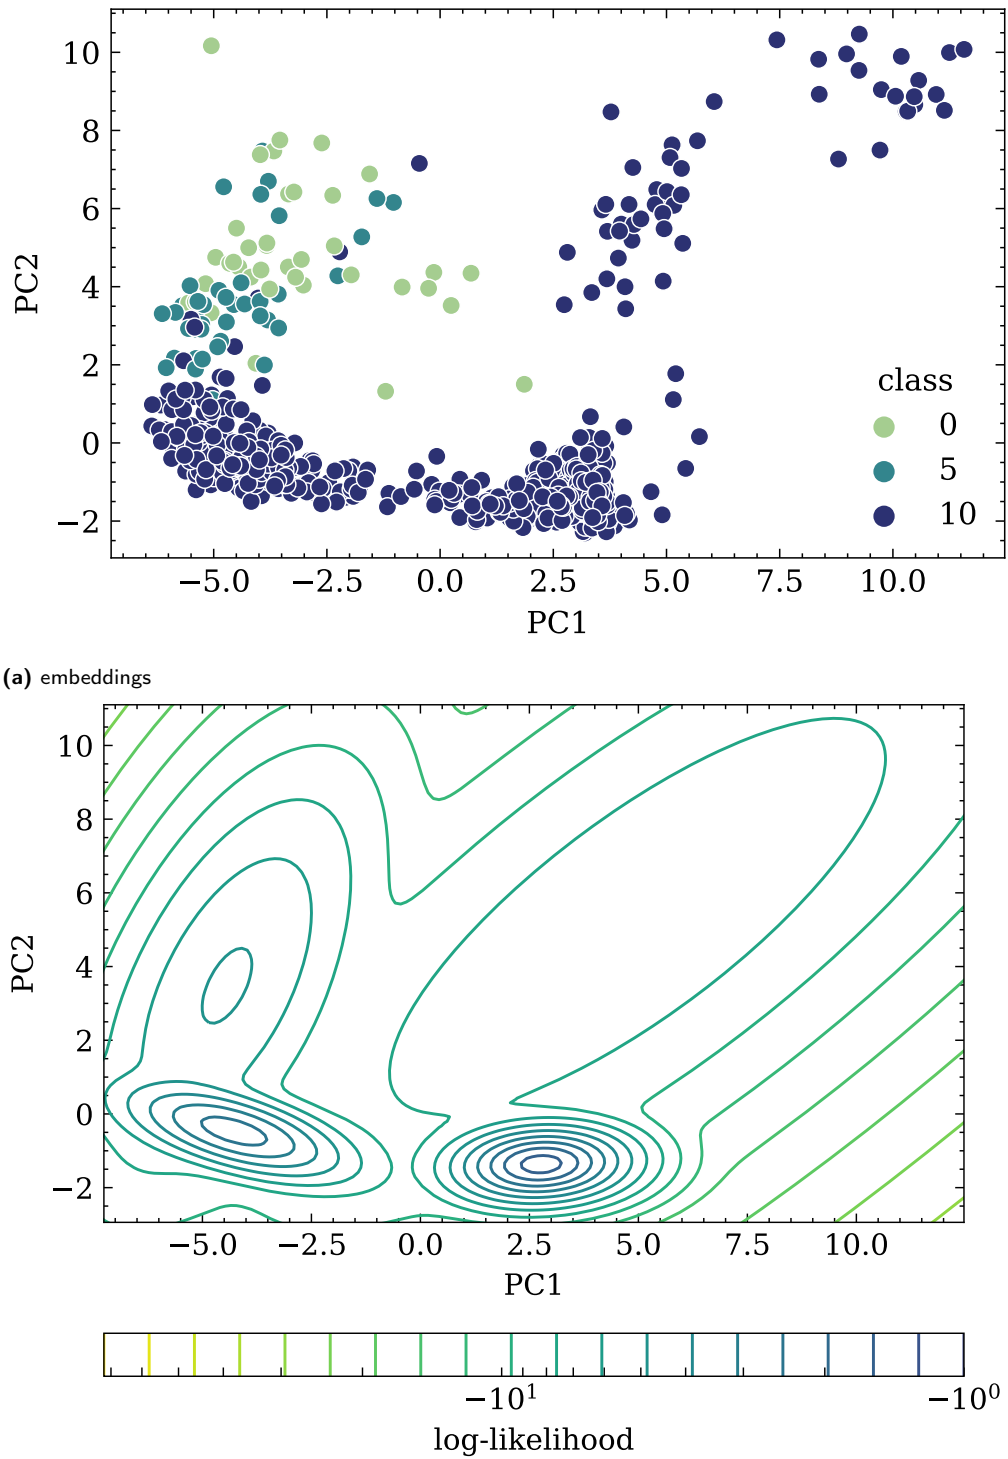
EEG embedding

Representing a typical 10-second EEG segment requires about 6.4×10^4 scalar values in it's raw form, thus special care is taken to overcome the curse of dimensionality. In this work, we embed the EEG data in a low-dimensional space, namely the space of GP parameters. We found the ictal embeddings to be distinguishable from the interictal embeddings even after the dimensionality reduction.

Following the described embedding of the EEG signal in the space of GP parameters, we adopt the new notation $\mathbb{E}_t \in \mathbb{R}^d$, where the new dimension satisfies $d \ll d' \times T$.

Gaussian process parameters embedding

Inference of Gaussian process (GP) parameters is a well-documented approach to modeling time-series data [Rasmussen, 2003]. The extension to multitask GPs enables modeling of multivariate time-series, such as



(b) log-likelihood as predicted by GMM

Figure 3.1. | Embedding of EEG. 2-channel segments of duration 10 (s), randomly sampled from the training set, are embedded in the GP parameter space. A GMM fit with expectation-maximization provides an approximate density estimation. Higher values are more likely to be generated by the GMM.

NS: fix class labels in legend, remove redundant axis label

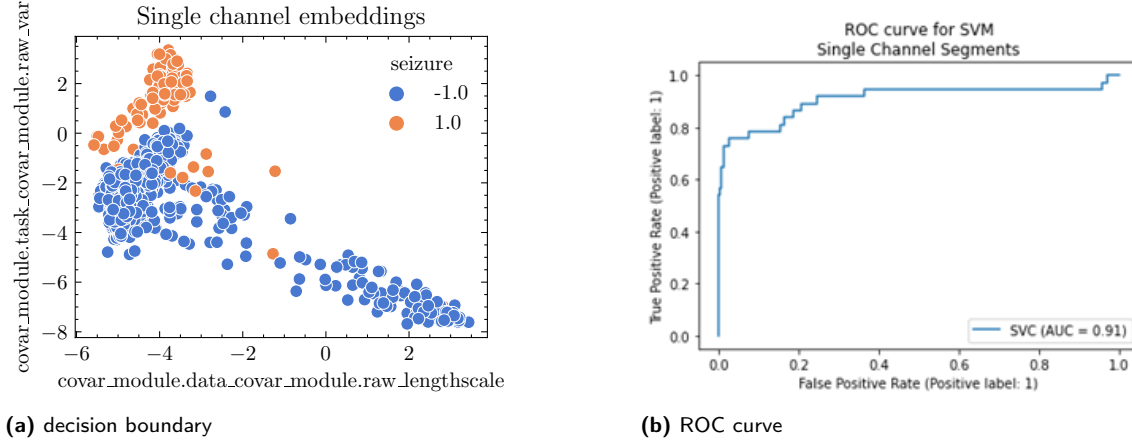


Figure 3.2. | Separability in the parameter space. A support vector machine achieves 0.91 AUC-ROC on a held-out test set for single-channel segments.

NS: remake seperability test figures

the case of multi-lead EEG signals.

For each EEG segment x , the preprocessing steps include:

- i Normalizing x by subtracting the mean and dividing by the standard deviation of the training set.
- ii initializing a GP model with zero mean, a scaled Matérn-1.5 kernel, and a rank-1 multitask covariance kernel.
- iii Optimizing the model's parameters to obtain a maximal marginal log-likelihood (details in appx. A).

The optimized model's parameters θ are persisted and used henceforth to represent the original EEG segment x .

$\mathcal{Z}(E)$ - Novelty Score

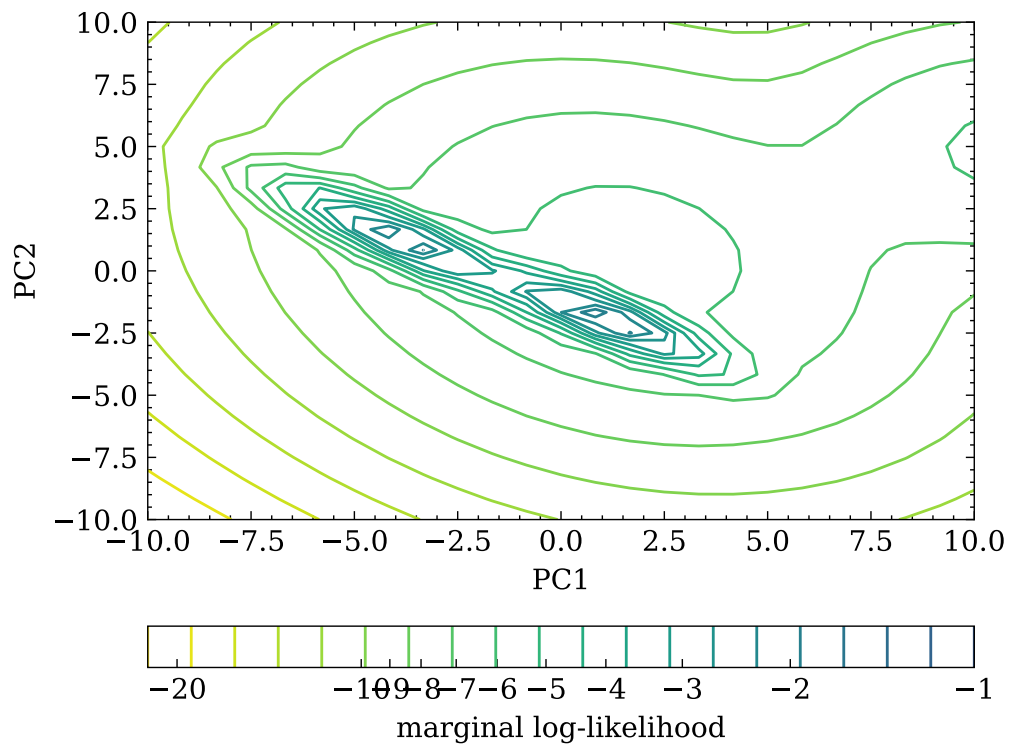
NS: reconsider if novelty score is necessary

We derive a generative, unsupervised novelty score $\mathcal{Z}(E)$ which will be used to define the seizure likelihood, $P(E | S) \propto \mathcal{Z}(E)$.


For a patient with epilepsy, we hypothesize that the anomalies in EEG data are inherently likely to reflect seizures. Thus, we substitute the likelihood term in Bayes' theorem with the novelty score $z(\theta)$ to provide an estimate of the likelihood of observing an EEG:

$$P(E | S) \equiv z(E) \quad (3.9)$$

We now show that our unsupervised anomaly detection method is suitable for detecting seizures in the dataset. Figure 3.5 shows interictal and ictal EEG samples, along with the calculated p-values for each clip.



(a) First and second principal components



(b) Pair plots for the first five principle components

Figure 3.3. | Gaussian Mixture Model Density Estimation. caption

NS: write caption




Figure 3.4. | Embedding EEG with GP hyperparameters. Posterior draws from a Gaussian process fit to maximize the likelihood of an observed single-channel EEG segment.

NS: replace with vectorized image

NS: add original sample

305 As can be seen, the interictal clip has a p-value of 0.4355 which indicates that it is in the normal distribution.
 306 On the other hand, the ictal clip has an extreme p-value of 0.9787 which indicates that it is a rare event, thus
 307 in a patient with epilepsy it is likely to be a seizure.

308 3.3. Weakly Supervised Bayesian Seizure Likelihood 309 Estimation

310 In the case of epilepsy, handling uncertainty in the face of evidence plays a major role, thus naturally appealing
 311 to the mathematical machinery termed Bayes' theorem².

312 We apply Bayes' theorem to estimate the updated likelihood of a seizure after observing an EEG signal:

$$313 \quad \mathbb{P}(S | E) = \frac{\mathbb{P}(E | S)\mathbb{P}(S)}{\mathbb{P}(E)} \propto \underbrace{\mathbb{P}(E | S)}_{\text{likelihood}} \cdot \underbrace{\mathbb{P}(S)}_{\text{prior}} \quad (3.10)$$

314 We now introduce our method to calculate the likelihood and prior, thus completing the description of
 315 our inference procedure.

$$316 \quad \mathbb{P}(S | E) = 1 - \frac{\mathbb{P}(\neg S)\mathbb{P}(E | \neg S)}{\mathbb{P}(E)} \quad (3.11)$$

317 due to the assumption $\mathbb{P}(S) + \mathbb{P}(\neg S) = 1$.

²much credit is due to Pierre Simon Laplace, who developed the form in common use today [McGrayne, 2011].

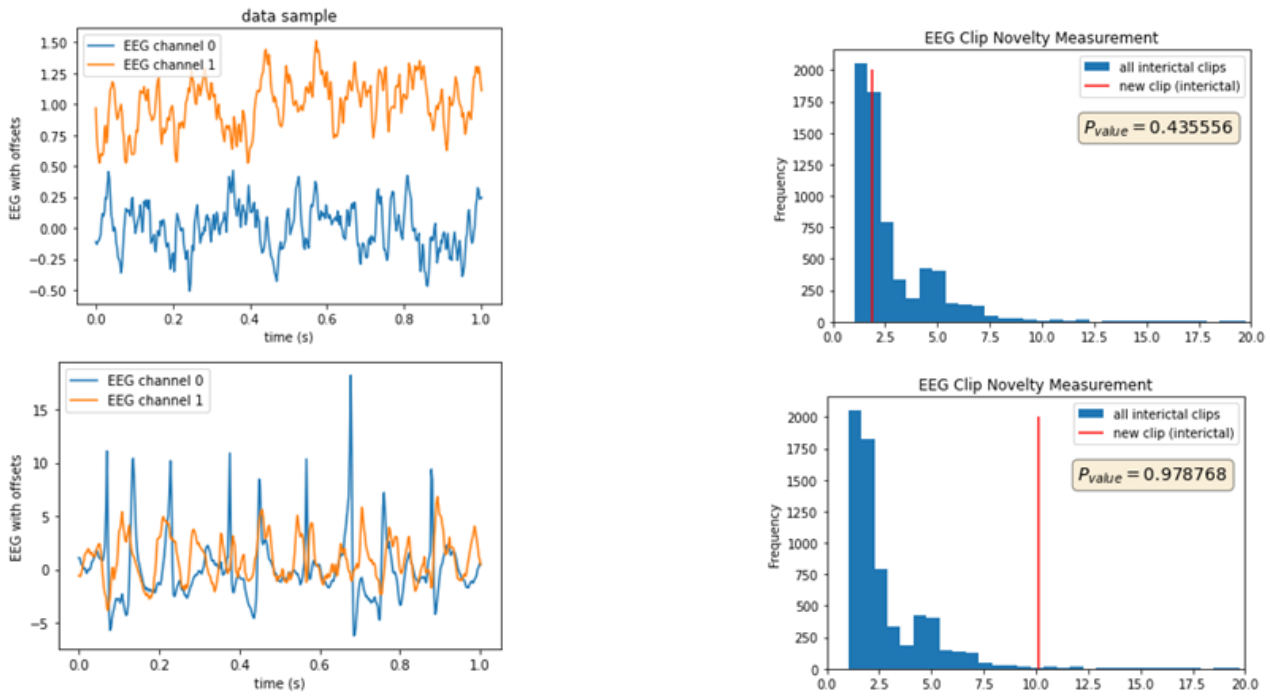


Figure 3.5. | Anomaly detection. An interictal double-channel segment has a p-value of 0.435 (top). An ictal double-channel segment has a p-value of 0.978 (bottom).

NS: remake example likelihood estimation figures

$$\mathbb{P}(E \mid \neg S) = \mathbb{P}(\{\hat{p}(e \mid \theta^*) \leq \hat{p}(E \mid \theta^*)\}) \quad (3.12)$$

$$\mathbb{P}(E) = \mathbb{P}(\{pdf(e) \leq pdf(E)\}) \quad (3.13)$$

3.3.1. Subject-dependent circadian-profile prior

Observing the circadian seizure distribution

NS: remake circadian profile diagrams panel horizontal

The von Mises (circular Gaussian) distribution

NS: format as definition

The von Mises distribution is defined as:




Figure 3.6. Circadian seizure distribution for I004_A0003_D001

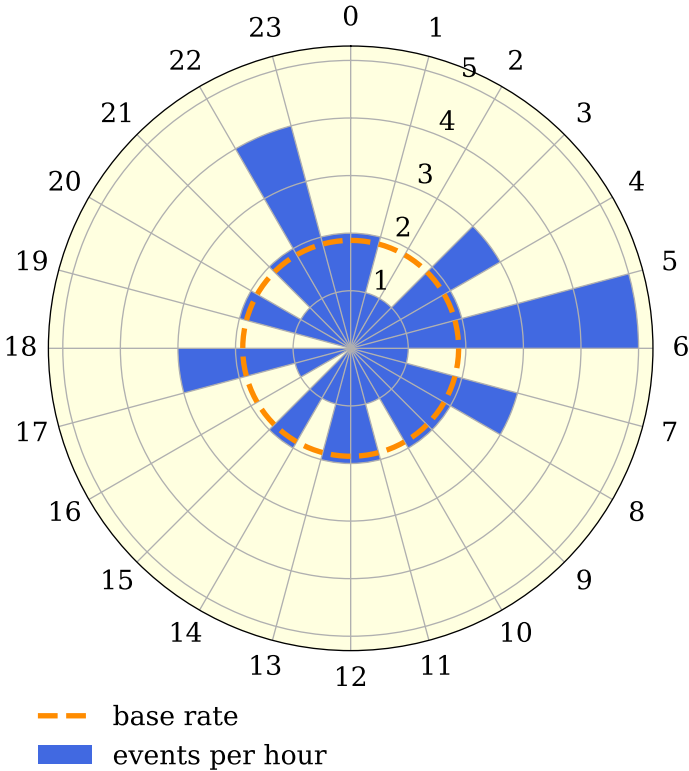


Figure 3.7. Circadian seizure distribution for I004_A0003_D001

$$f(x|\mu, \kappa; \omega) = \frac{\exp(\kappa \cos(\omega(x - \mu)))}{2\pi I_0(\kappa)} \quad (3.14)$$

Visually, the resulting distribution is similar to a bell-shaped normal distribution on a circle (see figure 3.8). In the definition above, μ determines the center of the bell, κ the spread, and ω scales the period length. In this work, we set $\omega \leftarrow \frac{2\pi}{24}$ to scale the period to 24-hours, and drop it from the notation for brevity in the following text.




Figure 3.8. | Von Mises distribution. The von Mises distribution probability density function (also known as the circular normal distribution) for $\mu = 8$, $k = 1/0.6$, $\omega = \frac{2\pi}{24}$

The *Circadian distribution prior* is the same one used by Karoly et al. [2017], based on a mixture of von Mises distributions (see equation 3.14).

Our work complements the current understanding of the prior’s contribution to predictive performance, and has many features that distinguish it from prior work.

NS: write list of ways this work differs from prior works

3.3.2. The Cox process prior

NS: write about Cox process (a.k.a. Inhomogeneous Poisson processes)

NS: make figure of inferring latent intensity

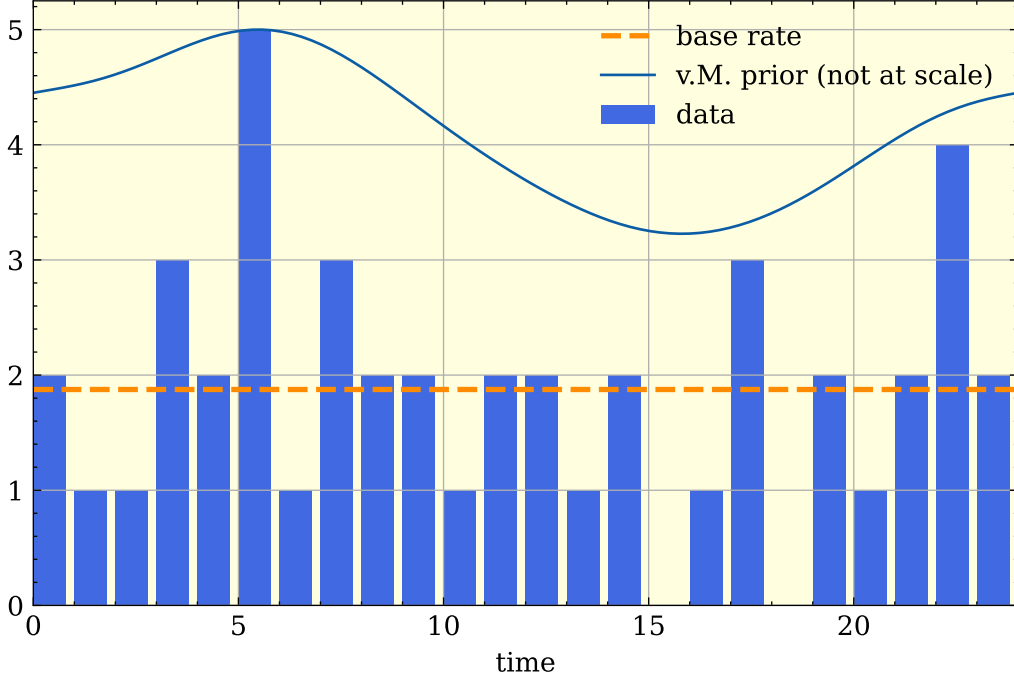


Figure 3.9. non-normalized v.M. prior overlaid on empirical circadian seizure distribution for I004_A0003_D001

3.3.3. Probabilistic model

We capture the notion of EEG observations and seizure events by random variables (r.v.s):

$$e_t \sim E(t) \in \Omega_E = \mathbb{R}^{c \times \tau} \quad (3.15)$$

$$s_t \sim S(t) \in \Omega_S = \{0, 1\} \quad (3.16)$$

where $\Omega_{E/S}$ is the sample space for each variable E and S , respectively. c is the number of EEG-channels, τ is the duration of EEG recorded (a.k.a. segment length). We define the event that a seizure occurred at time t as $\{S(t) = 1\}$.

Our work introduces a multilevel probabilistic model for dataset annotations (see figure 3.10). The model is inspired by the hypothesis that clinical annotations are very precise but not highly sensitive. That is, we believe that a proportion of seizures don't become recorded annotations, and explicitly take this belief into account in our model. These can often be explained by the annotator's reliance on video-footage and a possible lack of clinical symptom manifestations.

The model describes the process by which seizure annotations are generated. The W, k parameters control the shape of the individual's circadian profile. $\lambda(t)$ is the latent intensity function for the seizure-occurrence Cox process. After sampling a seizure-history $S(t)$, each seizure is dropped with a probability of p , and becomes an observable annotation with probability $1 - p$. This reflects annotators' missing sections of the EEG recordings, as well as sub-clinical seizures.

NS

add citation

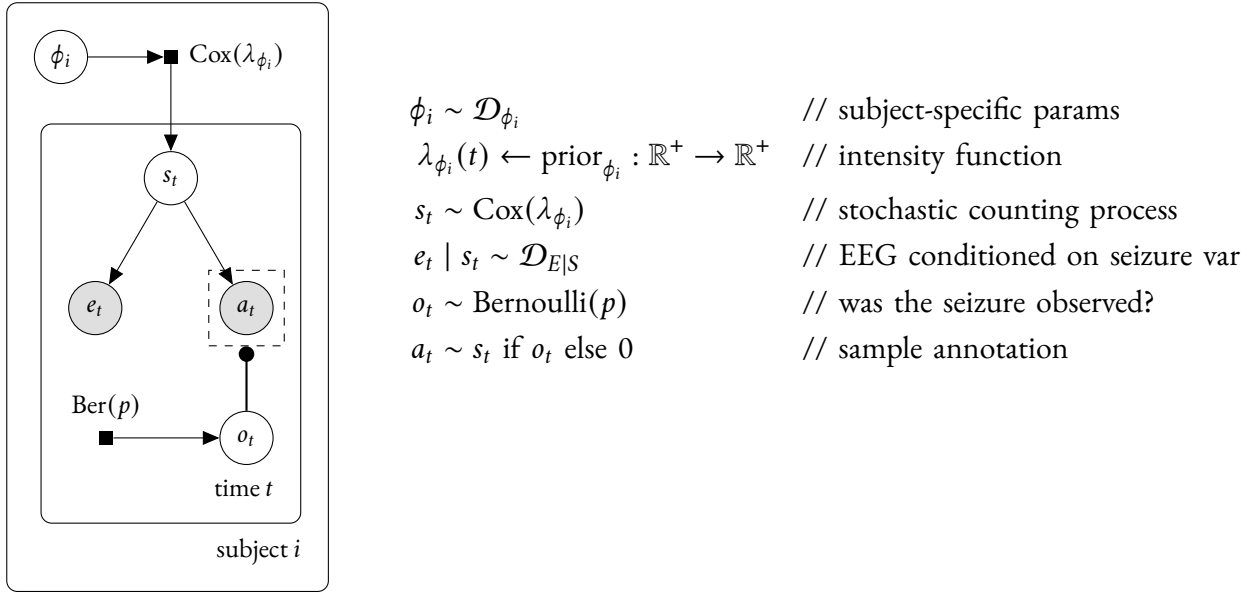


Figure 3.10. Probabilistically Modeling Seizure Occurrence, EEG signal and Annotations

356 More formally (see algorithm 1), the probability of obtaining an annotation at time t is a zero-inflated
 357 seizure occurrence model. Sequentially, the seizure-occurrence model is a Cox process with a stochastic
 358 intensity function $\lambda(t)$. In turn, the intensity function $\lambda(t)$ is a mixture model of 24 von-Mises distributions,
 359 one for each hour of the day. Finally, the weights W and common spread k parameters assume noninformative
 360 priors.

Algorithm 1 Seizure Annotation Model

$k \sim \text{Exp}(1)$
 $\alpha[0, \dots, 23] \leftarrow 1$
 $W[0, \dots, 23] \sim \text{Dir}(\alpha)$
 $\lambda(t) \leftarrow \sum_{i=0}^{23} W[i] \cdot f_{v.M.}(t; i, k)$
 $s_t \sim \text{Cox}(\lambda)$
 $a_t = \text{ZeroInflated}(s_t, p)$

Figure 3.11. A generative model for seizure events and annotations with a von-Mises-mixture intensity prior.

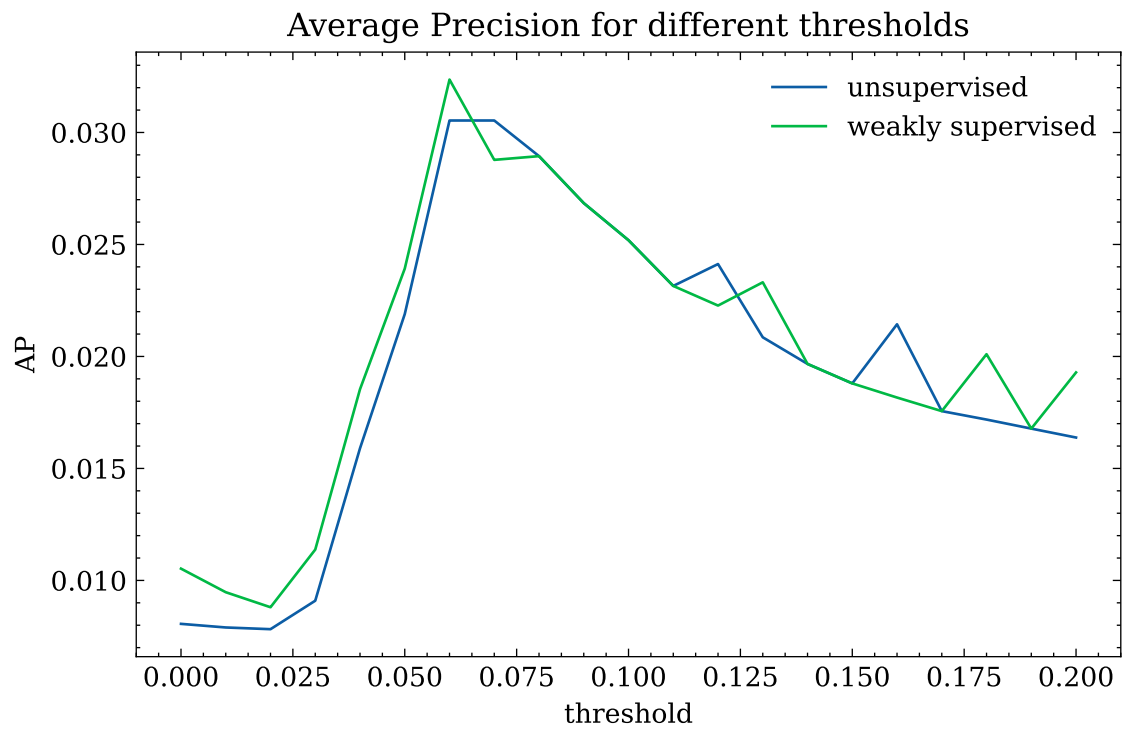


Figure 3.12. Average Precision scores (higher is better)

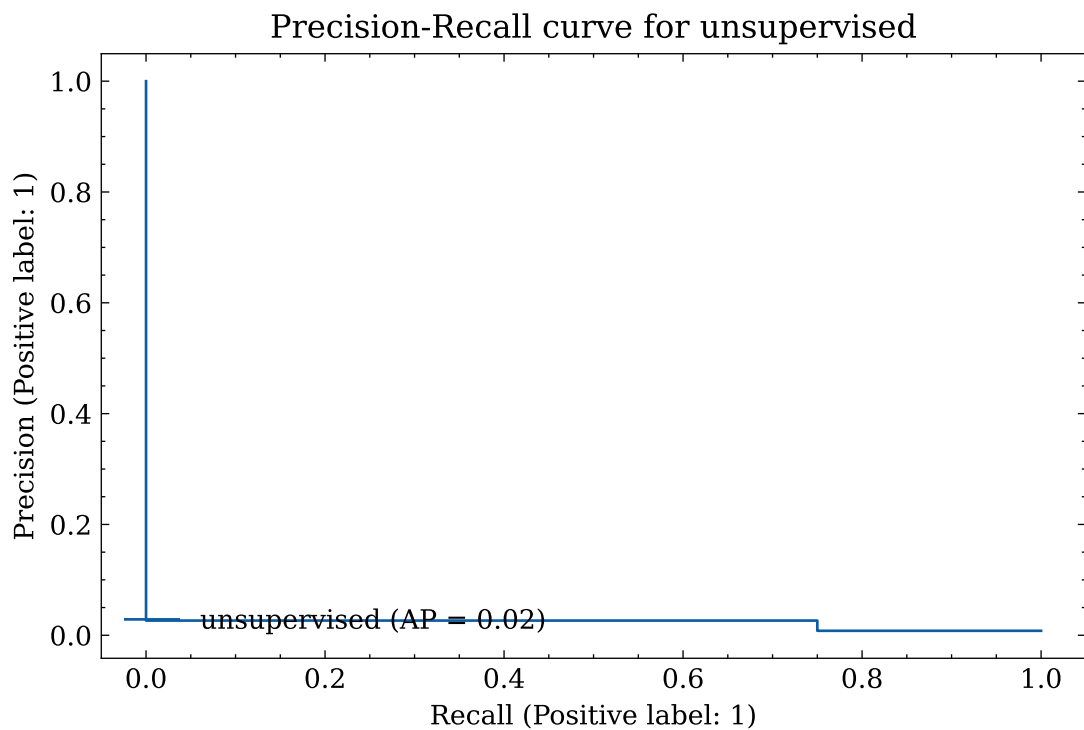


Figure 3.13. AP for unsupervised

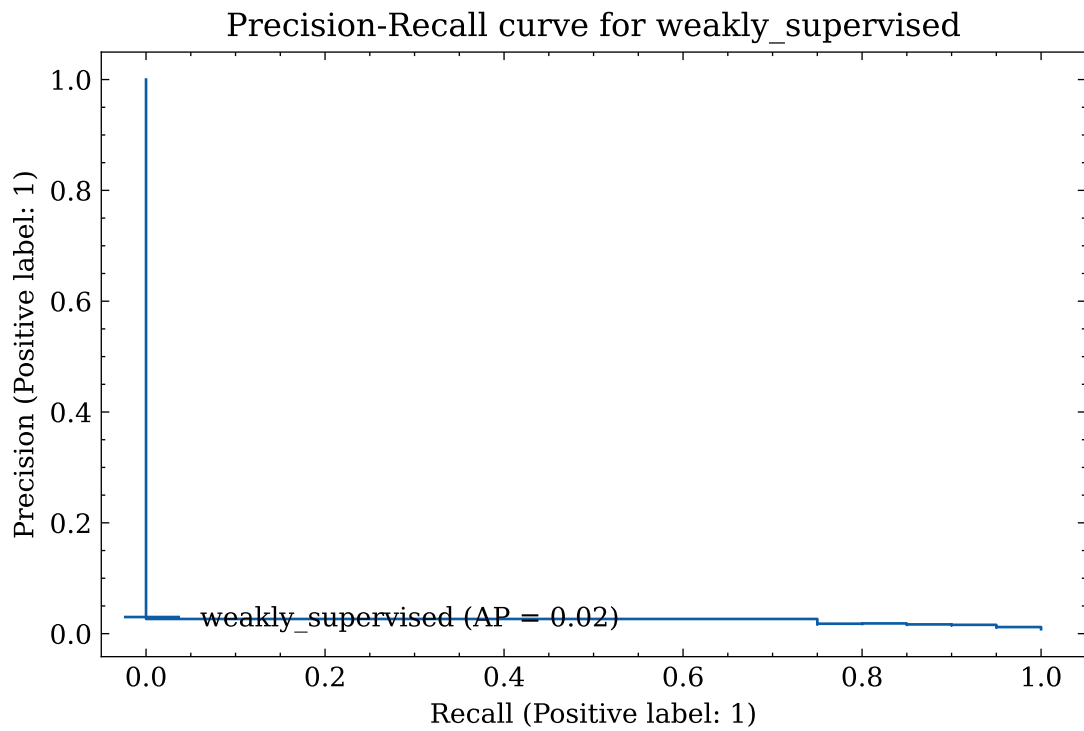


Figure 3.14. AP for weakly supervised (higher is better)

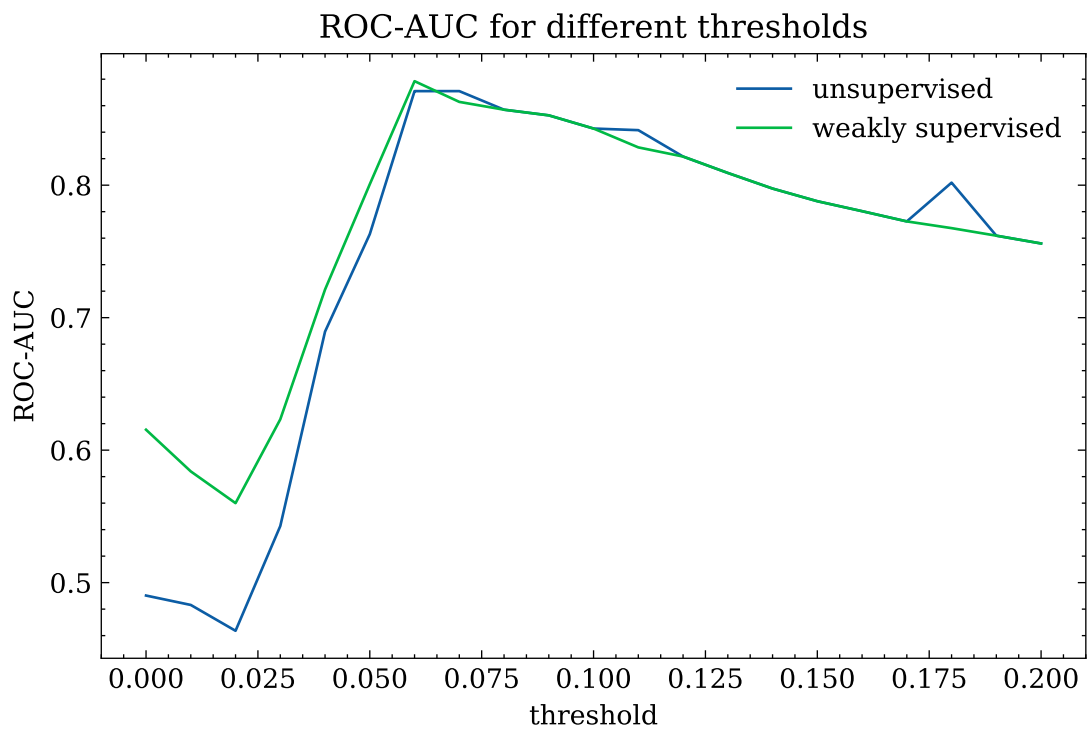


Figure 3.15. ROC-AUC scores (higher is better)

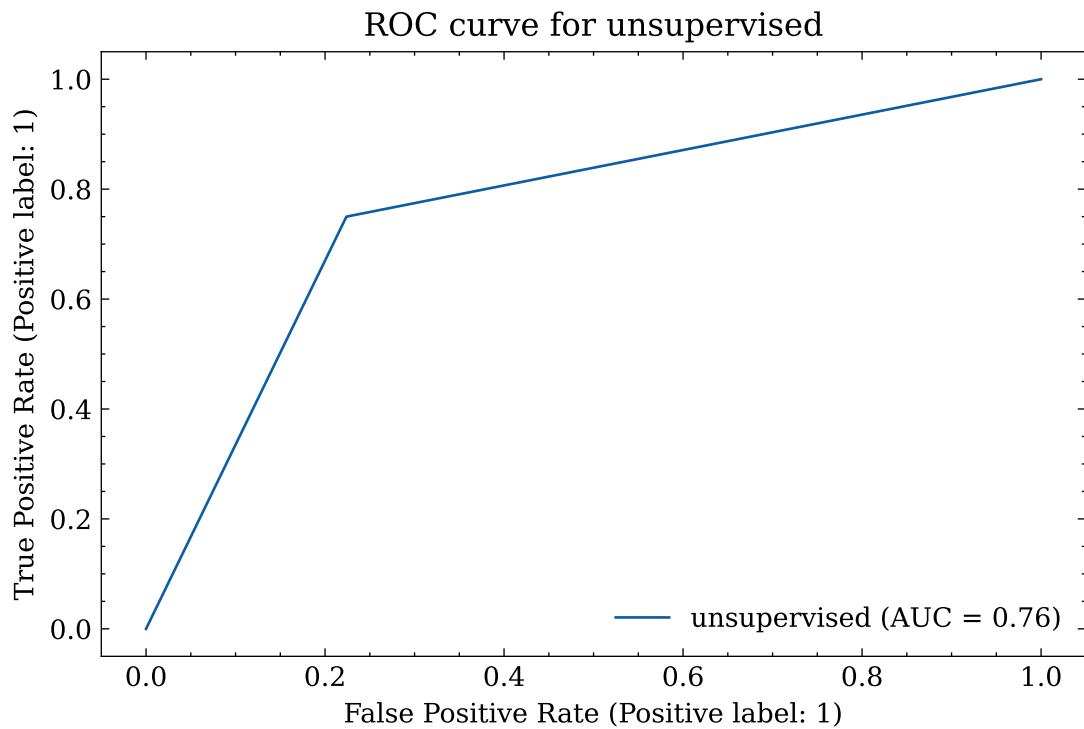


Figure 3.16. ROC for unsupervised




Figure 3.17. ROC for weakly supervised

NS: check which threshold made the roc curve fig and write it




Figure 3.18. Brier scores (lower is better)




Figure 3.19. AUC-ROC Heatmap for minutes time scale




Figure 3.20. AUC-ROC Heatmap for seconds time scale

3.3.4. Empirical Evaluation

Prior

Seizure Likelihood Estimation

Model validation

Linear probing

NS: rewrite Linear probing section

The *support vector machine* (SVM) is a popular method for solving problems in classification, regression and novelty detection [Bishop, 2006]. Fundamentally, it is a two-class linear classifier, i.e., it relies on linear models of the form

$$y(x) = w^T \phi(x) + b \quad (3.17)$$

where $\phi(x)$ denotes a fixed feature-space transformation, and b is the bias parameter. The training data set comprises N input vectors x_1, \dots, x_N , with corresponding target values t_1, \dots, t_N , where $t_n \in \{-1, 1\}$, and new data points x are classified according to the sign of $y(x)$.

374 By fitting the SVM to a training set and scoring the predictive accuracy on a hold-out test set, we can
375 quantify the linear separability of the dataset. This is used as a proxy for representation quality.

4. Conclusion

Historically, automatic seizure prediction based on EEG data has been researched since the 1970's, and continues to be an active field of research today [Mormann et al., 2007, Kuhlmann, Lehnertz, Richardson, Schelter, and Zaveri, 2018]. In the course of time, an attempt to characterize a *pre-ictal state* - a short-term (think minutes) pattern in neural activity, which upon manifestation is deterministically followed by a seizure - was dominantly pursued.

The notable methods reported in [Mirowski et al., 2009] influenced the field of seizure risk gauging. For example, 5 years later, the same correlation features were found in each of the three top-scoring submissions to the Kaggle seizure detection competition [?].

To this date, most published methods require labeled data, and suffer from high false alarm rates. Since they are fully-supervised, the cost dependence on expert-level domain knowledge makes these algorithms unsuitable for real-world forecasting devices, since unbiased expert labeling solutions are still not available. Furthermore, even clinical-grade annotations, considered to be the "gold standard" in seizure documentation, suffer from observation bias.

To overcome this problem, we propose a weakly-supervised Bayesian framework for likelihood estimation of epileptic seizures. We introduce a multilevel probabilistic model of seizure occurrences and clinical annotations (dataset labels). The model takes into account the possibility of missed seizures unrecorded by the annotator. We implemented an algorithm to compute the inference, namely Bayesian Seizure Likelihood Estimator (BSLE). Lastly, we validate the model and examine it's applicability to a practical seizure warning system.

NS

add citation

NS

add citation

4.1. Research questions

This thesis addresses the following research questions:

1. To what degree are various classifiers able to discriminate between interictal and ictal EEG segments, when using synchronicity-measures as features? (cf. chapter 2)
2. How can we model seizure likelihood estimation as a Bayesian inference problem, and how well does this model forecast seizures over different time-horizons? (cf. chapter 3)

4.2. Evaluation plan

Proper examination of seizure timing algorithms must account for all types of classification errors and be considerate of the imbalanced nature of the data. It is also essential to be able to benchmark the new

405 algorithm against existing methods. Perhaps above all, it is worthy to evaluate the algorithm's applicability to
406 real-world scenarios. Therefore, we will report a set of standard evaluation metrics for both the deterministic
407 and the probabilistic settings, and attempt to replicate the setting proposed in [Karoly et al., 2017] as much
408 as possible.

409 In evaluating the anomaly detection method proposed in this work, we will follow the guidelines provided
410 by the Kaggle Seizure Detection Challenge [UPenn and Mayo Clinic and Kaggle, 2014]. Namely, the model
411 will receive as input the training data from the Dog 1 dataset, and output class probabilities (0 for interictal,
412 1 for ictal). We will report ROC-curves and ROC-AUC on the held-out test set.

413 NS: report Brier Score

414 NS: report Snyder 2008

415 5. Discussion

416 Seizures are states of abnormal brain function, which are associated with unwanted symptoms such as
417 involuntary muscle movements (e.g., limb jerking, convulsing), partial or complete loss of consciousness and
418 memory relapses. Epilepsy is a condition in which unprovoked seizures occur recurrently. People suffering
419 from epilepsy may also suffer from side effects such as stress, fear, and sleep disorder. Prolonged seizures are
420 especially dangerous and can lead to sudden death. Moreover, the uncertainty regarding seizure occurrence
421 is regarded as one of the most disturbing factors of the disease.

422 5.1. Model validation

423 NS: write about model validation

424 NS: rewrite discussion. Points to include: (1) future work on hierarchical patient modeling

425 The problem of automatic seizure detection is challenging, inducing many attempts over the years. In
426 this work we attempted a probabilistic approach, relying on Bayes' rule to estimate the likelihood of a seizure
427 given an EEG segment. We assumed that seizures are rare, which led us to a novelty-score-based likelihood.
428 Following recent findings, we also assumed that seizures are approximately cyclical, taking this into account
429 in our prior.

430 Although the method we proposed works well, there are some drawbacks. First, the channel spatio-
431 temporal location is disregarded, and the pairing (in the double-channel case) was made arbitrarily. Further
432 work should utilize better channel selection and modeling the topographic qualities of the channels for
433 potentially improved results.

434 Second, the model of normal EEG was fit using the segments from the Canine-epilepsy-dataset which
435 were chosen for the Kaggle challenge. The class distribution differs from the true class distribution. To
436 combat this discrepancy, we dropped the ictal segments and used only the interictal segments for the model
437 of normality. Because ictal EEG is extremely less common than interictal EEG, we assume that using the
438 interictal segments is a sufficiently close approximation of the natural distribution. Further studies should
439 recalibrate the method based on the raw recordings in order to provide an even better estimate.

440 Thirdly, from the clinical perspective, the evaluation method is limited since the data originates from
441 canines with naturally occurring epilepsy, instead of humans. Dogs with naturally occurring epilepsy show
442 similar semiology to epilepsy in humans, but more work is required to validate this method on human EEG.

443 In summary, we present an anomaly detection method based on Gaussian processes embeddings, and
444 evaluate it on a seizure detection task. The method is significant because it is unsupervised, thus eliminated
445 the need for costly annotators.

446 We hope that future attempts at seizure likelihood estimation will utilize our Bayesian by incorporating
447 new likelihood functions and priors.

Bibliography

- Melinda S Kelley, Margaret P Jacobs, and Daniel H Lowenstein. The ninds epilepsy research benchmarks. *Epilepsia*, 50(3):579, 2009.
- Sonya B Dumanis, Jaqueline A French, Christophe Bernard, Gregory A Worrell, and Brandy E Fureman. Seizure forecasting from idea to reality. outcomes of the my seizure gauge epilepsy innovation institute workshop. *Eneuro*, 4(6), 2017.
- Piotr Mirowski, Deepak Madhavan, Yann LeCun, and Ruben Kuzniecky. Classification of patterns of eeg synchronization for seizure prediction. *Clinical neurophysiology*, 120(11):1927–1940, 2009.
- Eric R Kandel, James H Schwartz, Thomas M Jessell, Steven Siegelbaum, A James Hudspeth, and Sarah Mack. *Principles of neural science*, volume 4. McGraw-hill New York, 2000.
- Ingrid E Scheffer, Samuel Berkovic, Giuseppe Capovilla, Mary B Connolly, Jacqueline French, Laura Guilhoto, Edouard Hirsch, Satish Jain, Gary W Mathern, Solomon L Moshé, et al. Ilae classification of the epilepsies: position paper of the ilae commission for classification and terminology. *Epilepsia*, 58(4): 512–521, 2017.
- Palak Handa, Monika Mathur, and Nidhi Goel. Open and free eeg datasets for epilepsy diagnosis. *arXiv preprint arXiv:2108.01030*, 2021.
- Matthias Ihle, Hinnerk Feldwisch-Drentrup, César A Teixeira, Adrien Witon, Björn Schelter, Jens Timmer, and Andreas Schulze-Bonhage. Epilepsiae—a european epilepsy database. *Computer methods and programs in biomedicine*, 106(3):127–138, 2012.
- J-B Schiratti, Jean-Eudes Le Douget, Michel Le van Quyen, Slim Essid, and Alexandre Gramfort. An ensemble learning approach to detect epileptic seizures from long intracranial eeg recordings. In *2018 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 856–860. IEEE, 2018.
- Klaus Lehnertz, Christian Geier, Thorsten Rings, and Kirsten Stahn. Capturing time-varying brain dynamics. *EPJ Nonlinear Biomedical Physics*, 5:2, 2017.
- Michael G Rosenblum, Arkady S Pikovsky, and Jürgen Kurths. From phase to lag synchronization in coupled chaotic oscillators. *Physical Review Letters*, 78(22):4193, 1997.
- Jean-Philippe Lachaux, Eugenio Rodriguez, Jacques Martinerie, and Francisco J Varela. Measuring phase synchrony in brain signals. *Human brain mapping*, 8(4):194–208, 1999.

- 477 Florian Mormann, Ralph G Andrzejak, Christian E Elger, and Klaus Lehnertz. Seizure prediction: the long
478 and winding road. *Brain*, 130(2):314–333, 2007.
- 479 Piotr W Mirowski, Yann LeCun, Deepak Madhavan, and Ruben Kuzniecky. Comparing svm and convolu-
480 tional networks for epileptic seizure prediction from intracranial eeg. In *2008 IEEE workshop on machine*
481 *learning for signal processing*, pages 244–249. IEEE, 2008.
- 482 Lars Buitinck, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad
483 Niculae, Peter Prettenhofer, Alexandre Gramfort, Jaques Grobler, Robert Layton, Jake VanderPlas,
484 Arnaud Joly, Brian Holt, and Gaël Varoquaux. API design for machine learning software: experiences
485 from the scikit-learn project. In *ECML PKDD Workshop: Languages for Data Mining and Machine*
486 *Learning*, pages 108–122, 2013.
- 487 Philippa J Karoly, Hoameng Ung, David B Grayden, Levin Kuhlmann, Kent Leyde, Mark J Cook, and
488 Dean R Freestone. The circadian profile of epilepsy improves seizure forecasting. *Brain*, 140(8):2169–
489 2182, 2017.
- 490 Maxime O Baud, Jonathan K Kleen, Emily A Mirro, Jason C Andrechak, David King-Stephens, Edward F
491 Chang, and Vikram R Rao. Multi-day rhythms modulate seizure risk in epilepsy. *Nature communications*,
492 9(1):1–10, 2018.
- 493 Maxime O Baud, Timothée Proix, Vikram R Rao, and Kaspar Schindler. Chance and risk in epilepsy.
494 *Current opinion in neurology*, 33(2):163–172, 2020.
- 495 CE Rasmussen. Cki williams gaussian processes for machine learning, 2006.
- 496 Sergios Theodoridis. *Machine learning: a Bayesian and optimization perspective*. Academic press, 2015.
- 497 Kathryn A Davis, Beverly K Sturges, Charles H Vite, Vanessa Ruedebusch, Gregory Worrell, Andrew B
498 Gardner, Kent Leyde, W Douglas Sheffield, and Brian Litt. A novel implanted device to wirelessly record
499 and analyze continuous intracranial canine eeg. *Epilepsy research*, 96(1-2):116–122, 2011.
- 500 UPenn and Mayo Clinic and Kaggle. seizure detection & prediction challenges, 2014.
501 <https://www.kaggle.com/c/seizure-detection> and [https://www.kaggle.com/c/](https://www.kaggle.com/c/seizure-prediction)
502 [seizure-prediction](https://www.kaggle.com/c/seizure-prediction).
- 503 Carl Edward Rasmussen. Gaussian processes in machine learning. In *Summer school on machine learning*,
504 pages 63–71. Springer, 2003.
- 505 Sharon Bertsch McGrayne. The theory that would not die. In *The Theory That Would Not Die*. Yale
506 University Press, 2011.
- 507 Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*.
508 Springer-Verlag, Berlin, Heidelberg, 2006. ISBN 0387310738.
- 509 Levin Kuhlmann, Klaus Lehnertz, Mark P Richardson, Björn Schelter, and Hitten P Zaveri. Seizure
510 prediction—ready for a new era. *Nature Reviews Neurology*, 14(10):618–630, 2018.

- 511 Jacob R Gardner, Geoff Pleiss, David Bindel, Kilian Q Weinberger, and Andrew Gordon Wilson. Gpy-
512 torch: Blackbox matrix-matrix gaussian process inference with gpu acceleration. In *Advances in Neural*
513 *Information Processing Systems*, 2018.
- 514 William Falcon and The PyTorch Lightning team. PyTorch Lightning, 3 2019. URL [https://github.](https://github.com/PyTorchLightning/pytorch-lightning)
515 [com/PyTorchLightning/pytorch-lightning](https://github.com/PyTorchLightning/pytorch-lightning).

Appendices

517 A. GP embedding: training details

518 For dimensionality reduction of an observed EEG segment, we performed exact inference of the GP pa-
 519 rameters maximizing the observation likelihood, using the GPyTorch and PyTorch Lightning frameworks
 520 [Gardner et al., 2018, Falcon and The PyTorch Lightning team, 2019].

521 More formally, in fitting the Gaussian processes to the EEG samples we carried out exact Type-II Maximum
 522 Likelihood Estimation for each sample. This means optimizing the model’s hyperparameters (mean module,
 523 covariance module, etc.) w.r.t maximization of the *marginal log likelihood* (MLL) of the given data \mathbf{E}, \mathbf{t} :

$$524 \quad \theta \leftarrow \operatorname{argmax}_{\theta} p_f(\mathbf{E} | \mathbf{t}) = \int p(\mathbf{E} | f(\mathbf{t}))p(f(\mathbf{t}) | \mathbf{t})df \quad (\text{A.1})$$

525 where $f \sim \mathcal{GP}(\mu, K)$ is the modeled signal before adding the homoscedastic Gaussian noise, and θ is the
 526 set of parameters to be optimized. See code for implementation details.

GP & inference (training) configuration details		
	Parameter	Value
GP params	mean module	zero mean
	covariance module	scaled Matérn-1.5 kernel
	task covariance rank	1
	number of tasks	2
	likelihood (noise model)	Gaussian (homoscedastic)
Training params	optimizer	Adam
	learning rate	0.01
	max. number of epochs	1000
	patience (early stopping)	8

Table A.1. The GP parameters and training parameters used in our experiments.

תקציר

התקציר שלי

חיזוי הסתברותי של התקפי אפילפסיה

מחקר לשם מילוי חלקי של הדרישות לקבלת תואר
"מגיסטר למדעים"

מגיש

נועם סיגל

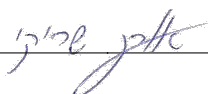
העבודה נעשתה בהדרכת

ד"ר אורן שריקי

וד"ר דוד טולפין

המחלקה למדעי המחשב

אישור המנחים:



אישור להשלים:

אב ה'תשפ"ב

באר שבע

חיזוי הסתברותי של התקפי אפילפסיה

מחקר לשם מילוי חלקי של הדרישות לקבלת תואר
"מגיסטר למדעים"

מגיש
נועם סיגל



הוגש למחלקה למדעי המחשב באוניברסיטת בן גוריון בנגב

אב ה'תשפ"ב

באר שבע